

# Optimization of neural networks

Dmitry Yarotsky

# Parametrized predictive models

True response function:  $y = f(\mathbf{x})$ , where  $\mathbf{x}$  is the input vector

- $y \in \mathbb{R}$  for regression
- $y \in \{0, 1\}$  for binary classification

Predictive model:  $y = \tilde{f}(\mathbf{x}, \mathbf{W})$ , and  $\mathbf{W}$  are model parameters (e.g., network weights)

“Soft classification”:  $\tilde{f}(\mathbf{x}, \mathbf{W}) \in [0, 1]$

# Model training as a parametric optimization

Loss function:  $L(\mathbf{W}) = \int l(f(\mathbf{x}), \tilde{f}(\mathbf{x}, \mathbf{W})) d\mu(\mathbf{x})$

Sample average measure  $\mu$  :

- $d\mu(\mathbf{x}) = \frac{1}{N} \sum_{n=1}^N \delta(\mathbf{x} - \mathbf{x}_n)$  with Dirac's delta — “finite training set” scenario
- $d\mu(\mathbf{x}) = p(\mathbf{x}) d\mathbf{x}$  with some (e.g. Gaussian) density  $p(\mathbf{x})$  — “population average” scenario

Function  $l(f(\mathbf{x}), \tilde{f}(\mathbf{x}, \mathbf{W}))$  measures the discrepancy between  $f$  and  $\tilde{f}$ , e.g.:

- Regression:  $l(y, \tilde{y}) = \frac{1}{2}(y - \tilde{y})^2$
- Classification:  $l(y, \tilde{y}) = -y \log \tilde{y} - (1 - y) \log(1 - \tilde{y})$

Model training:

$$L(\mathbf{W}) \longrightarrow \min_{\mathbf{W}}$$

# Gradient-based optimization

- $\mathbf{W}$  high-dimensional
- $L(\mathbf{W})$  non-smooth, non-convex

Most popular approach: gradient-based optimization and its modifications

Basic gradient descent with learning rate  $\alpha \in (0, 1)$ :

$$\mathbf{W}^{(n+1)} = \mathbf{W}^{(n)} - \alpha \nabla_{\mathbf{W}} L(\mathbf{W}^{(n)})$$

Gradient descent with Nesterov momentum ( $\alpha, \beta \in (0, 1)$ ):

$$\mathbf{W}^{(n+1)} = \mathbf{W}^{(n)} - \mathbf{V}^{(n)}$$

$$\mathbf{V}^{(n+1)} = \alpha \nabla_{\mathbf{W}} L(\mathbf{W}^{(n)}) + \beta \mathbf{V}^{(n)}$$

**Exercise:** how can we interpret the coefficients  $\alpha$  and  $\beta$ ?

## Computation of $\nabla_{\mathbf{W}}L$ : “Error backpropagation”

$$\nabla_{\mathbf{W}}L(\mathbf{W}) = \int \frac{\partial l}{\partial \tilde{y}}(f(\mathbf{x}), \tilde{f}(\mathbf{x}, \mathbf{W})) \cdot \nabla_{\mathbf{W}}\tilde{f}(\mathbf{x}, \mathbf{W}) d\mu(\mathbf{x})$$

$$\nabla_{\mathbf{W}}\tilde{f} = (\nabla_{\mathbf{w}_1}\tilde{f}, \dots, \nabla_{\mathbf{w}_K}\tilde{f})$$

$\frac{\partial l}{\partial \tilde{y}}(f(\mathbf{x}), \tilde{f}(\mathbf{x}, \mathbf{W}))$ : directly computed from  $y$  and  $\tilde{y}$


$\tilde{y} = \tilde{f}(\mathbf{x}, \mathbf{W})$ : “forward propagation”

To find  $\nabla_{\mathbf{W}}\tilde{f}$ : use layerwise representation

$$\tilde{f}(\mathbf{x}, \mathbf{W}) = g_K(g_{K-1}(\dots g_1(\mathbf{x}, \mathbf{w}_1), \dots \mathbf{w}_{K-1}), \mathbf{w}_K)$$

$\mathbf{z}_k$ : output of the  $k$ 'th layer (known from “forward propagation”)

$$\mathbf{z}_k = g_k(\mathbf{z}_{k-1}, \mathbf{w}_k)$$



# “Error backpropagation”



$$\nabla_{\mathbf{w}_K} \tilde{f}(\mathbf{x}, \mathbf{W}) = \frac{\partial g_K}{\partial \mathbf{w}_K}(\mathbf{z}_{K-1}, \mathbf{w}_K)$$

$$\begin{aligned} \nabla_{\mathbf{w}_{K-1}} \tilde{f}(\mathbf{x}, \mathbf{W}) &= \nabla_{\mathbf{w}_K} g_K(g_{K-1}(\mathbf{z}_{K-2}, \mathbf{w}_{K-1}), \mathbf{w}_K) \\ &= \frac{\partial g_K}{\partial \mathbf{z}_{K-1}}(\mathbf{z}_{K-1}, \mathbf{w}_K) \cdot \frac{\partial g_{K-1}}{\partial \mathbf{w}_{K-1}}(\mathbf{z}_{K-2}, \mathbf{w}_{K-1}) \end{aligned}$$

...

$$\begin{aligned} \nabla_{\mathbf{w}_k} \tilde{f}(\mathbf{x}, \mathbf{W}) &= \frac{\partial g_K}{\partial \mathbf{z}_{K-1}}(\mathbf{z}_{K-1}, \mathbf{w}_K) \cdots \frac{\partial g_{k+1}}{\partial \mathbf{z}_k}(\mathbf{z}_k, \mathbf{w}_{k+1}) \\ &\quad \cdot \frac{\partial g_k}{\partial \mathbf{w}_k}(\mathbf{z}_{k-1}, \mathbf{w}_k) \end{aligned}$$

# Basic theory: convergence of gradient descent

See e.g. Yu. Nesterov, Introductory Lectures on Convex Programming Volume I: Basic course.

- **Global minima:**  $L(\mathbf{W}_*) = \min_{\mathbf{W} \in \mathbb{R}^W} L(\mathbf{W})$
- **Local minima:**  $L(\mathbf{W}_*) = \min_{\mathbf{W} \in U} L(\mathbf{W})$ , where  $U$  is an open neighborhood of  $\mathbf{W}_*$
- **Stationary points:**  $\nabla_{\mathbf{W}} L(\mathbf{W}_*) = 0$  (assuming  $L(\mathbf{W})$  is smooth)

In general, gradient descent converges to a stationary point.

## Proposition

*Suppose function  $L$  is lower bounded and  $L \in \mathcal{W}^{2,\infty}(\mathbb{R}^W)$ , so that  $|\nabla L(\mathbf{a}) - \nabla L(\mathbf{b})| \leq M|\mathbf{a} - \mathbf{b}|$  with some Lipschitz constant  $M$ . Let  $\alpha < \frac{2}{M}$ . Then  $\nabla L(\mathbf{W}^{(n)}) \rightarrow 0$ , and  $\min_{n=1,\dots,N} |\nabla L(\mathbf{W}^{(n)})| = O(N^{-1/2})$ .*

# Proof

$$\begin{aligned} L(\mathbf{W}^{(n+1)}) &= L(\mathbf{W}^{(n)}) + \left\langle \mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}, \int_0^1 \nabla L(\mathbf{W}^{(n)} + t(\mathbf{W}^{(n+1)} - \mathbf{W}^{(n)})) dt \right\rangle \\ &\leq L(\mathbf{W}^{(n)}) + \langle \mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}, \nabla L(\mathbf{W}^{(n)}) \rangle \\ &\quad + |\mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}| \int_0^1 Mt |\mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}| dt \\ &\leq L(\mathbf{W}^{(n)}) + \langle \mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}, \nabla L(\mathbf{W}^{(n)}) \rangle + \frac{M}{2} |\mathbf{W}^{(n+1)} - \mathbf{W}^{(n)}|^2 \\ &\leq L(\mathbf{W}^{(n)}) + (-\alpha + \frac{M}{2} \alpha^2) |\nabla L(\mathbf{W}^{(n)})|^2 \\ &< L(\mathbf{W}^{(n)}) \end{aligned}$$

if  $\alpha < \frac{2}{M}$ . Let  $c = \alpha(1 - \frac{M}{2}\alpha) > 0$ , then

$$\sum_{n=1}^N |\nabla L(\mathbf{W}^{(n)})|^2 \leq \frac{1}{c} \sum_{n=1}^N (L(\mathbf{W}^{(n)}) - L(\mathbf{W}^{(n+1)})) \leq \frac{1}{c} (L(\mathbf{W}^{(1)}) - \min_{\mathbf{W}} L(\mathbf{W})),$$

$$\min_{n=1, \dots, N} |\nabla L(\mathbf{W}^{(n)})| \leq \frac{c^{-1/2}}{\sqrt{N}} (L(\mathbf{W}^{(1)}) - \min_{\mathbf{W}} L(\mathbf{W}))^{1/2}$$



## Formulation in terms of stopping condition

Assume the **stopping condition**:  $|\nabla L(\mathbf{W}^{(n)})| < \epsilon$ .

Then, optimization terminates in  $O\left(\frac{L(\mathbf{W}^{(1)}) - \min_{\mathbf{W}} L(\mathbf{W})}{\epsilon^2}\right)$  steps.

**Exercise:** What is the optimal value of  $\alpha$ , assuming  $M$  is known?

**Exercise:** Give an example of gradient descent converging to a stationary point which is not a (local or global) minimum.

# Linearization and spectral analysis

Suppose that  $L \in C^2(\mathbb{R})$  and  $\mathbf{W}_*$  is a stationary point.  
For  $\mathbf{W}$  near  $\mathbf{W}_*$ :

$$\nabla L(\mathbf{W}) = D^2L(\mathbf{W}_*) \cdot (\mathbf{W} - \mathbf{W}_*) + o(|\mathbf{W} - \mathbf{W}_*|),$$

where  $D^2L(\mathbf{W}_*)$  is the Hessian matrix. Optimization iterates:

$$\begin{aligned}\mathbf{W}^{(n+1)} - \mathbf{W}_* &= \mathbf{W}^{(n)} - \mathbf{W}_* - \alpha \nabla L(\mathbf{W}^{(n)}) \\ &= \mathbf{W}^{(n)} - \mathbf{W}_* - \alpha D^2L(\mathbf{W}_*) \cdot (\mathbf{W}^{(n)} - \mathbf{W}_*) + o(|\mathbf{W}^{(n)} - \mathbf{W}_*|) \\ &= (1 - \alpha D^2L(\mathbf{W}_*)) \cdot (\mathbf{W}^{(n)} - \mathbf{W}_*) + o(|\mathbf{W}^{(n)} - \mathbf{W}_*|)\end{aligned}$$



Convergence is determined by eigenvalues of  $D^2L(\mathbf{W}_*)$ :

- positive: convergence
- negative: divergence

# Evasion of saddle points

**Saddle points:**  $D^2L(\mathbf{W}_*)$  has both positive and negative eigenvalues

Typically, saddles are evaded by optimization, due to the presence of diverging components in  $\mathbf{W}^{(n)} - \mathbf{W}_*$ . The manifold of converging  $\mathbf{W}^{(n)}$  has Lebesgue measure 0.<sup>1</sup>




Chi Jin and M. Jordan, How to Escape Saddle Points Efficiently: Saddle points can slow down optimization; perturbing the GD can help.

---

<sup>1</sup>B. Recht, Saddles Again

- No smoothness, in general (e.g. with ReLU): local minima of  $L(\mathbf{W})$  are non-differentiable



Th. Laurent, J. von Brecht, The  
Multilinear Structure of ReLU  
Networks, arXiv:1712.10132

- Large size of the network and its structure are important

# Empirical observations of real-life ANNs

From A. Choromanska et al., The Loss Surfaces of Multilayer Networks, arXiv:1412.0233:

- Large networks train well despite their size. Optimization can terminate at different local minima, but they seem to be equivalent and yield similar performance on a test set.
- The probability of finding a bad (high value) local minimum is non-zero for small-size networks and decreases quickly with network size.
- Struggling to find the global minimum on the training set (as opposed to one of the many good local ones) is not useful in practice and may lead to overfitting.

# Conceptual pictures of the loss surface (conjectured)

From M. Baity-Jesi, *Comparing Dynamics: Deep Neural Networks versus Glassy Systems*: two alternatives

- ① The loss landscape is very rough, has many isolated local minima, but GD tends to find good minima having low loss.
- ② The loss function is highly nonlinear, but has few local minima, and the minima are connected. (Example:  
$$L(w_1, w_2) = (w_2 - w_1^2)^2 + \epsilon w_1^2.$$
)




(Note: these conceptual pictures have only a limited value due to the “curse of dimensionality” in  $\mathbb{R}^W$ , lack of characterization of locality and depth of a local minimum, etc.)

## Some current research directions

- Numerical studies of loss surface and gradient descent
- Direct analytic studies of simple (toy) scenarios:
  - Deep linear networks (no nonlinear activation)
  - Wide shallow networks with small training sets, pyramidal networks (no spurious local minima)
- Large-size limits:
  - Large-width limit: Gaussian approximation for signal propagation, connections to random matrices and spherical spin glasses
  - Phenomenological: Stochastic PDE and Langevin dynamics
- Specialized networks (e.g., convnets)

# F. Draxler et al., Essentially No Barriers in Neural Network Energy Landscape, arXiv:1803.00885

- Two local minima are connected by a path, and then it is deformed to find a low loss trajectory
- ResNets and DenseNets on CIFAR10 and CIFAR100
- The optimized path: approximately constant loss



(Global description of the optimal set?)



# I. Safran, O. Shamir, Spurious Local Minima are Common in Two-Layer ReLU Neural Networks, arXiv:1712.08968

**Spurious local minimum**  $\mathbf{W}_0$ :  $\min_{\mathbf{W}} L(\mathbf{W}) < L(\mathbf{W}_0) < L(\mathbf{W}')$  for  $\mathbf{W}'$  in a small neighborhood of  $\mathbf{W}_0$

## Theorem

*Consider the optimization problem*

$$\min_{\mathbf{w}_1, \dots, \mathbf{w}_n \in \mathbb{R}^k} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mathbf{0}, I)} \left( \sum_{i=1}^n (\mathbf{w}_i^\top \mathbf{x})_+ - \sum_{i=1}^k (\mathbf{v}_i^\top \mathbf{x})_+ \right)^2,$$

*where  $\mathbf{v}_1, \dots, \mathbf{v}_k$  are orthogonal unit vectors in  $\mathbb{R}^k$ . Then for  $n = k \in \{6, 7, \dots, 20\}$  as well as  $(k, n) \in \{(8, 9), (10, 11), \dots, (19, 20)\}$ , this objective function has spurious local minima.*

Proof: computer-assisted

## Dependence on $n, k$

More spurious minima observed at larger  $k$ , but overparametrization (large  $n$ ) appears to partly remove them.

**Table:** Spurious local minima found for  $n = k$

| k  | n  | % of runs converging to local minima | Average minimal eigenvalue | Average objective value |
|----|----|--------------------------------------|----------------------------|-------------------------|
| 6  | 6  | 0.3%                                 | 0.0047                     | 0.025                   |
| 7  | 7  | 5.5%                                 | 0.014                      | 0.023                   |
| 8  | 8  | 12.6%                                | 0.021                      | 0.021                   |
| 9  | 9  | 21.8%                                | 0.027                      | 0.02                    |
| 10 | 10 | 34.6%                                | 0.03                       | 0.022                   |
| 11 | 11 | 45.5%                                | 0.034                      | 0.022                   |
| 12 | 12 | 58.5%                                | 0.035                      | 0.021                   |
| 13 | 13 | 73%                                  | 0.037                      | 0.022                   |
| 14 | 14 | 73.6%                                | 0.038                      | 0.023                   |
| 15 | 15 | 80.3%                                | 0.038                      | 0.024                   |
| 16 | 16 | 85.1%                                | 0.038                      | 0.027                   |
| 17 | 17 | 89.7%                                | 0.039                      | 0.027                   |
| 18 | 18 | 90%                                  | 0.039                      | 0.029                   |
| 19 | 19 | 93.4%                                | 0.038                      | 0.031                   |
| 20 | 20 | 94%                                  | 0.038                      | 0.033                   |

**Table:** Spurious local minima found for  $n \neq k$

| k  | n  | % of runs converging to local minima | Average minimal eigenvalue | Average objective value |
|----|----|--------------------------------------|----------------------------|-------------------------|
| 8  | 9  | 0.1%                                 | 0.0059                     | 0.021                   |
| 10 | 11 | 0.1%                                 | 0.0057                     | 0.018                   |
| 11 | 12 | 0.1%                                 | 0.0056                     | 0.017                   |
| 12 | 13 | 0.3%                                 | 0.0054                     | 0.016                   |
| 13 | 14 | 1.5%                                 | 0.0015                     | 0.038                   |
| 14 | 15 | 5.5%                                 | 0.002                      | 0.033                   |
| 15 | 16 | 10.1%                                | 0.004                      | 0.032                   |
| 16 | 17 | 18%                                  | 0.0055                     | 0.031                   |
| 17 | 18 | 20.9%                                | 0.007                      | 0.031                   |
| 18 | 19 | 36.9%                                | 0.0064                     | 0.028                   |
| 19 | 20 | 49.1%                                | 0.0077                     | 0.027                   |